

## In this notebook

Many college courses conclude by giving students the opportunity to evaluate the course and the instructor anonymously. In the article “Beauty in the Classroom: Professors’ Pulchritude and Putative Pedagogical Productivity” ([PDF](#)), authors Daniel Hamermesh and Amy M. Parker suggest (based on a data set of teaching evaluation scores collected at UT Austin) that student evaluation scores can partially be predicted by features unrelated to teaching, such as the physical attractiveness of the instructor.

In this notebook, we will use this data to try and predict a course- and instructor-specific “baseline” score (excluding the effect of teaching quality), against which to measure instructor performance.

## Attribution

Parts of this lab are based on a lab assignment from the OpenIntro textbook “Introductory Statistics with Randomization and Simulation” that is released under a Creative Commons Attribution-ShareAlike 3.0 Unported license. The book website is at <https://www.openintro.org/book/isrs/>.

## Data

The data were gathered from end of semester student evaluations for a large sample of professors from the University of Texas at Austin. In addition, six students looked at a photograph of each professor in the sample, and rated the professors’ physical appearance. More specifically:

Each of the professors’ pictures was rated by each of six undergraduate students: Three women and three men, with one of each gender being a lower-division, two upper-division students (to accord with the distribution of classes across the two levels). The raters were told to use a 10 (highest) to 1 rating scale, to concentrate on the physiognomy of the professor in the picture, to make their ratings independent of age, and to keep 5 in mind as an average.

We are using a slightly modified version of the original data set from the published paper. The dataset was released along with the textbook “Data Analysis Using Regression and Multilevel/Hierarchical Models” (Gelman and Hill, 2007).

## Setup

We will start by importing relevant libraries, setting up our notebook, reading in the data, and checking that it was loaded correctly.

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn import metrics
from sklearn import model_selection
from sklearn.linear_model import LinearRegression

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
```

```
url = 'https://www.openintro.org/stat/data/evals.csv'
df = pd.read_csv(url)
df.head()
df.columns
df.shape
```

```

    score      rank      ethnicity gender language age  cls_perc_eval \
0    4.7  tenure track      minority female  english   36    55.81395
1    4.1  tenure track      minority female  english   36    68.80000
2    3.9  tenure track      minority female  english   36    60.80000
3    4.8  tenure track      minority female  english   36    62.60163
4    4.6      tenured  not minority    male  english   59    85.00000

    cls_did_eval  cls_students  cls_level  ...  cls_credits  bty_follower \
0             24             43    upper  ...  multi credit          5
1             86            125    upper  ...  multi credit          5
2             76            125    upper  ...  multi credit          5
3             77            123    upper  ...  multi credit          5
4             17             20    upper  ...  multi credit          4

    bty_fiupper  bty_f2upper  bty_m1lower  bty_m1upper  bty_m2upper  bty_avg \
0              7              6              2              4              6      5.0
1              7              6              2              4              6      5.0
2              7              6              2              4              6      5.0
3              7              6              2              4              6      5.0
4              4              2              2              3              3      3.0

    pic_outfit  pic_color
0  not formal    color
1  not formal    color
2  not formal    color
3  not formal    color
4  not formal    color

[5 rows x 21 columns]

```

```

Index(['score', 'rank', 'ethnicity', 'gender', 'language', 'age',
      'cls_perc_eval', 'cls_did_eval', 'cls_students', 'cls_level',
      'cls_profs', 'cls_credits', 'bty_follower', 'bty_fiupper', 'bty_f2upper',
      'bty_m1lower', 'bty_m1upper', 'bty_m2upper', 'bty_avg', 'pic_outfit',
      'pic_color'],
      dtype='object')

```

```
(463, 21)
```

Each row in the data frame represents a different course, and columns represent features of the courses and professors. Here's the data dictionary:

variable	description
score	average professor evaluation score: (1) very unsatisfactory - (5) excellent.
rank	rank of professor: teaching, tenure track, tenured.
ethnicity	ethnicity of professor: not minority, minority.
gender	gender of professor: female, male.
language	language of school where professor received education: english or non-english.
age	age of professor.
cls_perc_eval	percent of students in class who completed evaluation.
cls_did_eval	number of students in class who completed evaluation.
cls_students	total number of students in class.

variable	description
cls_level	class level: lower, upper.
cls_profs	number of professors teaching sections in course in sample: single, multiple.
cls_credits	number of credits of class: one credit (lab, PE, etc.), multi credit.
btty_f1lower	beauty rating of professor from lower level female: (1) lowest - (10) highest.
btty_f1upper	beauty rating of professor from upper level female: (1) lowest - (10) highest.
btty_f2upper	beauty rating of professor from second upper level female: (1) lowest - (10) highest.
btty_m1lower	beauty rating of professor from lower level male: (1) lowest - (10) highest.
btty_m1upper	beauty rating of professor from upper level male: (1) lowest - (10) highest.
btty_m2upper	beauty rating of professor from second upper level male: (1) lowest - (10) highest.
btty_avg	average beauty rating of professor.
pic_outfit	outfit of professor in picture: not formal, formal.
pic_color	color of professor's picture: color, black & white.

Source: [OpenIntro book](#).

Note that:

- `score` is the target variable - this is what we want our model to predict. We expect that the score is a function of the teaching quality, characteristics of the course, and non-teaching related characteristics of the instructor. However, the “true” teaching quality for each course is not known.
- the variables that begin with a `cls_` prefix are features that relate to the course. These features could potentially affect student evaluations: for example, students may rank one-credit lab courses more highly than multi-credit lecture courses.
- variables such as `rank`, `ethnicity`, `gender`, `language`, `age`, and the variables with a `btty_` prefix are features that relate to the instructor. They do not necessarily to the quality of instruction! These features may also affect student evaluations: for example, students may rate instructors more highly if they are physically attractive.
- variables with the `pic_` prefix describe the photograph that was shown to the students who provided the `btty_` scores. This should have no effect on the student evaluations, since those were evaluations by students who were enrolled in the course (not the students who were shown the photograph and asked to provide an attractiveness score.) (For your reference: on the bottom of page 7 of the paper, the authors describe why they include this variable and how they used it )

## Explore data

As always, start by exploring the data:

```
df.describe()
```

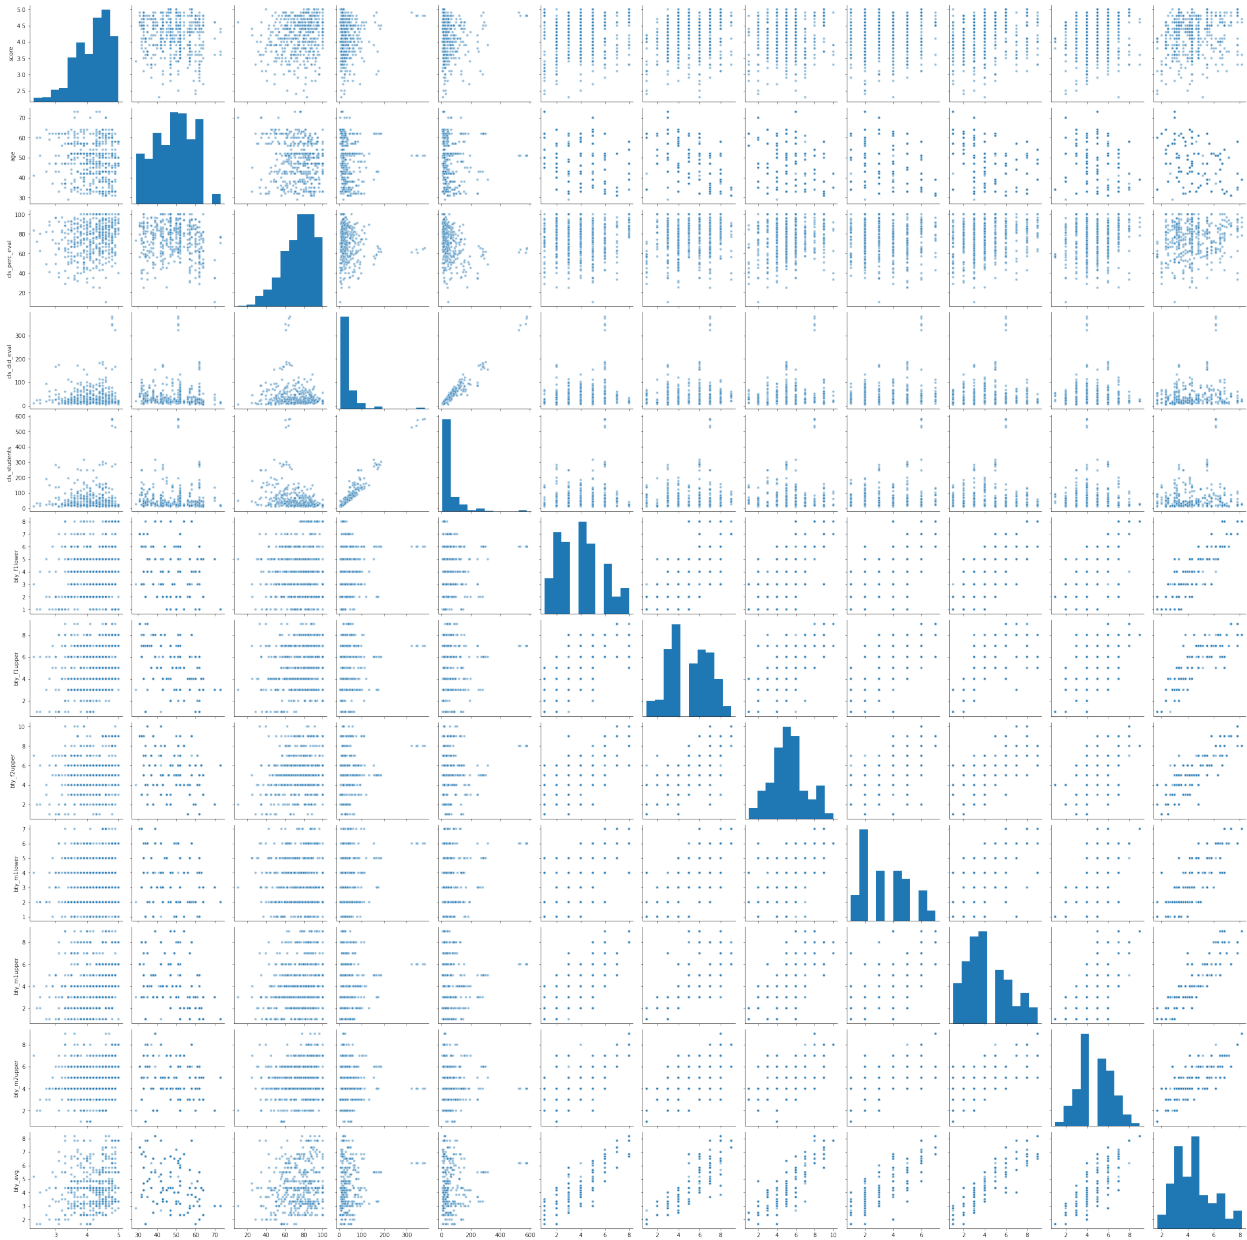
	score	age	cls_perc_eval	cls_did_eval	cls_students	\
count	463.000000	463.000000	463.000000	463.000000	463.000000	
mean	4.174730	48.365011	74.427788	36.624190	55.177106	
std	0.543865	9.802742	16.756311	45.018481	75.072800	
min	2.300000	29.000000	10.416670	5.000000	8.000000	
25%	3.800000	42.000000	62.696165	15.000000	19.000000	
50%	4.300000	48.000000	76.923080	23.000000	29.000000	
75%	4.600000	57.000000	87.249170	40.000000	60.000000	
max	5.000000	73.000000	100.000000	380.000000	581.000000	
	btty_f1lower	btty_f1upper	btty_f2upper	btty_m1lower	btty_m1upper	\
count	463.000000	463.000000	463.000000	463.000000	463.000000	

mean	3.963283	5.019438	5.213823	3.412527	4.146868
std	1.873936	1.934437	2.018224	1.637102	2.110586
min	1.000000	1.000000	1.000000	1.000000	1.000000
25%	2.000000	4.000000	4.000000	2.000000	3.000000
50%	4.000000	5.000000	5.000000	3.000000	4.000000
75%	5.000000	7.000000	6.000000	5.000000	5.000000
max	8.000000	9.000000	10.000000	7.000000	9.000000

	bty_m2upper	bty_avg
count	463.000000	463.000000
mean	4.751620	4.417844
std	1.575266	1.527380
min	1.000000	1.667000
25%	4.000000	3.167000
50%	5.000000	4.333000
75%	6.000000	5.500000
max	9.000000	8.167000

```
sns.pairplot(df, plot_kws={'alpha':0.5, 'size': 0.1})
```

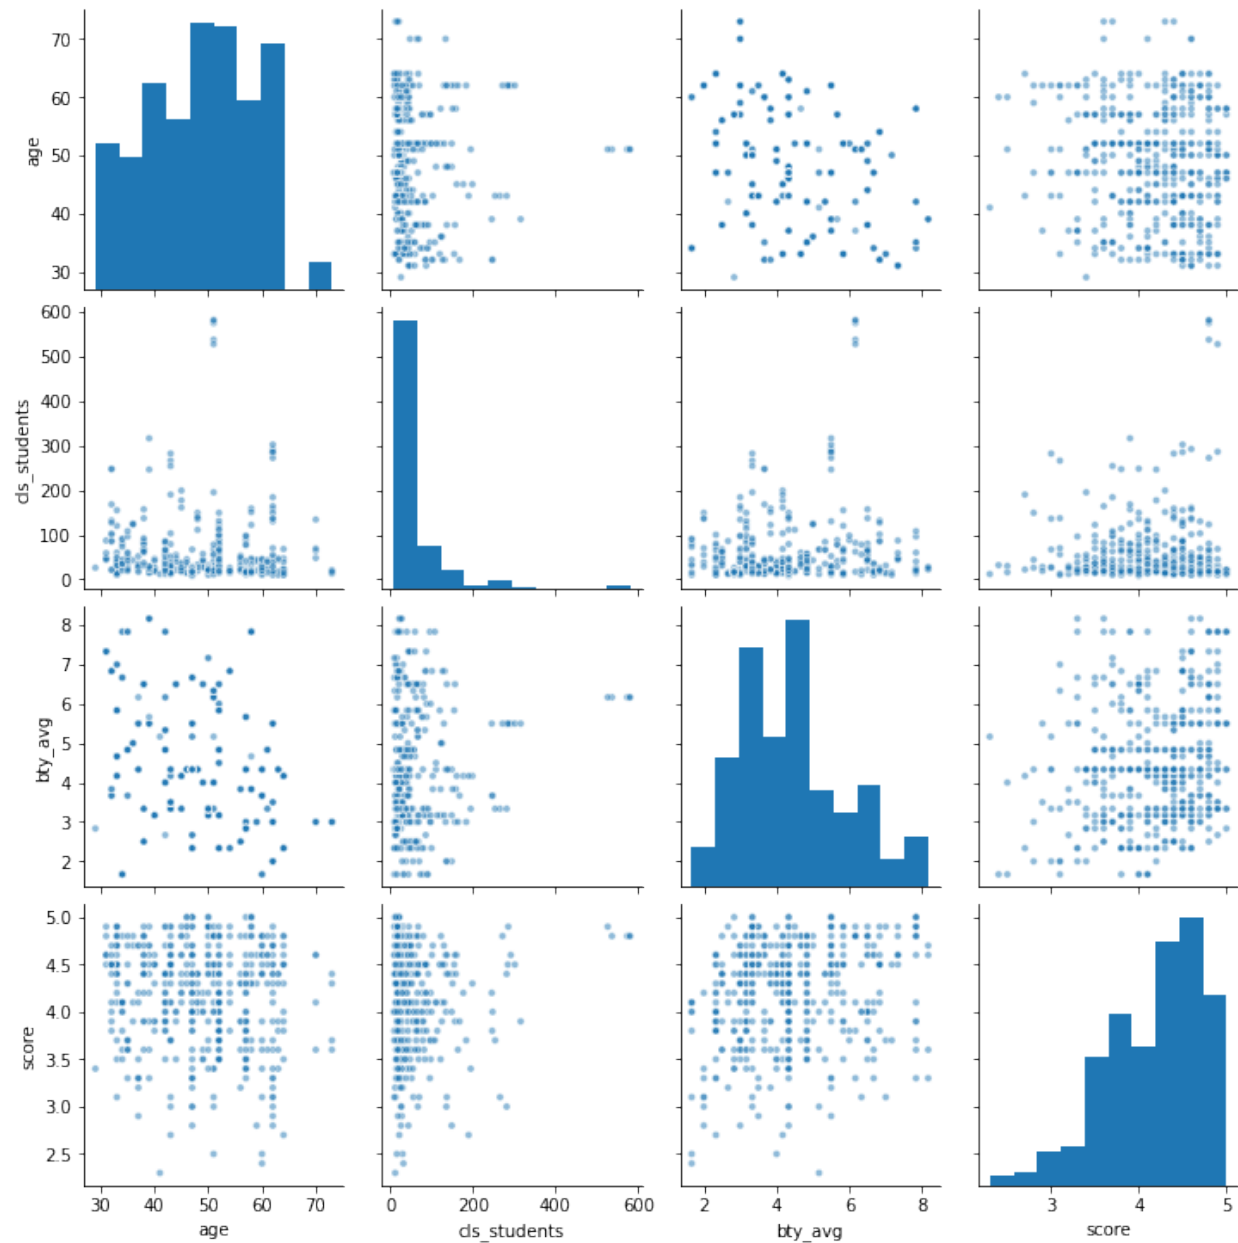
```
<seaborn.axisgrid.PairGrid at 0x7f76a4dd44f0>
```



With so many numeric variables, the pair plot is hard to read. We can create a pairplot excluding some variables that we don't expect to be useful for visualization: `cls_perc_eval`, `cls_did_eval`. We will also exclude the individual attractiveness ratings `btv_f1lower`, `btv_f1upper`, `btv_f2upper`, `btv_m1lower`, `btv_m1upper`, `btv_m2upper`, since the overall attractiveness rating is still represented by `btv_avg`.

```
sns.pairplot(df, vars=['age', 'cls_students', 'btv_avg', 'score'], plot_kws={'alpha':0.5,
'size': 0.1})
```

```
<seaborn.axisgrid.PairGrid at 0x7f76a01a3430>
```



As part of our exploration of the data, we can also examine the effect of non-numeric variables related to the instructor and the class: rank, ethnicity, gender, language, cls\_level, cls\_profs, cls\_credits.

```
for feature in ['rank', 'ethnicity', 'gender', 'language', 'cls_level', 'cls_profs',
               'cls_credits']:
    df.groupby([feature])['score'].describe()
```

	count	mean	std	min	25%	50%	75%	max
rank								
teaching	102.0	4.284314	0.498263	3.3	3.9	4.40	4.7	5.0
tenure track	108.0	4.154630	0.561104	2.3	3.7	4.35	4.6	4.9
tenured	253.0	4.139130	0.550262	2.4	3.8	4.20	4.6	5.0

ethnicity									
minority	64.0	4.071875	0.581588	2.7	3.675	4.05	4.525	5.0	
not minority	399.0	4.191228	0.536505	2.3	3.850	4.30	4.600	5.0	

	count	mean	std	min	25%	50%	75%	max	
gender									
female	195.0	4.092821	0.563814	2.3	3.7	4.1	4.5	5.0	
male	268.0	4.234328	0.521896	2.4	3.9	4.3	4.6	5.0	

	count	mean	std	min	25%	50%	75%	max	
language									
english	435.0	4.189655	0.547183	2.3	3.9	4.30	4.6	5.0	
non-english	28.0	3.942857	0.434979	3.4	3.6	3.75	4.4	4.8	

	count	mean	std	min	25%	50%	75%	max	
cls_level									
lower	157.0	4.238217	0.592532	2.5	3.8	4.4	4.7	5.0	
upper	306.0	4.142157	0.515104	2.3	3.8	4.2	4.5	5.0	

	count	mean	std	min	25%	50%	75%	max	
cls_profs									
multiple	306.0	4.184641	0.551177	2.4	3.8	4.3	4.6	5.0	
single	157.0	4.155414	0.530529	2.3	3.8	4.3	4.6	5.0	

	count	mean	std	min	25%	50%	75%	max	
cls_credits									
multi credit	436.0	4.147018	0.542464	2.3	3.8	4.2	4.6	5.0	
one credit	27.0	4.622222	0.334357	3.5	4.5	4.7	4.9	5.0	

**Discussion Question 1** Describe the relationship between `score` and the overall attractiveness rating `btv_avg`. Is there an apparent correlation? If so, is it a positive or a negative correlation? What about `age` and `cls_students`, do they appear to be correlated with `score`?

Also describe the relationship between `score` and the categorical variables you explored above that are related to characteristics of the *instructor*: `rank`, `ethnicity`, `gender`, `language`. Which of these variables have an apparent correlation with `score`? Is it a positive or a negative correlation?

Are any of the apparent relationships you observed unexpected to you? Explain.

## Encoding categorical variables

To represent a categorical variable (with no inherent ordering) in a regression, we can use *one hot encoding*. It works as follows:

- For a categorical variable  $x$  with values  $1, \dots, M$
- Represent with  $M$  binary features:  $\phi_1, \phi_2, \dots, \phi_M$
- Model in regression  $w_1\phi_1 + \dots + w_M\phi_M$

We can use the `get_dummies` function in `pandas` for one hot encoding. Create a copy of the dataframe with all categorical variables transformed into indicator (“dummy”) variables, and save it in a new data frame called `df_enc`.

Compare the columns of the `df` data frame versus the `df_enc` data frame.

```
df_enc = pd.get_dummies(df)
df_enc.columns
```

```
Index(['score', 'age', 'cls_perc_eval', 'cls_did_eval', 'cls_students',
      'bty_f1lower', 'bty_f1upper', 'bty_f2upper', 'bty_m1lower',
      'bty_m1upper', 'bty_m2upper', 'bty_avg', 'rank_teaching',
      'rank_tenure track', 'rank_tenured', 'ethnicity_minority',
      'ethnicity_not minority', 'gender_female', 'gender_male',
      'language_english', 'language_non-english', 'cls_level_lower',
      'cls_level_upper', 'cls_profs_multiple', 'cls_profs_single',
      'cls_credits_multi credit', 'cls_credits_one credit',
      'pic_outfit_formal', 'pic_outfit_not formal', 'pic_color_black&white',
      'pic_color_color'],
      dtype='object')
```

## Split data

Next, we split the encoded data into a training set (70%) and test set (30%). We will be especially interested in evaluating the model performance on the test set. Since it was not used to train the model parameters (intercept and coefficients), the performance on this data gives us a better idea of how the model may perform on new data.

We'll use the `train_test_split` method in `sklearn's model_selection` module. Since it randomly splits the data, we'll pass a random "state" into the function that makes the split repeatable (same split every time we run this notebook) and ensures that everyone in the class will have exactly the same split.

```
train, test = model_selection.train_test_split(df_enc, test_size=0.3, random_state=9)
# why 9? see https://dilbert.com/strip/2001-10-25
train.shape
test.shape
```

```
(324, 31)
```

```
(139, 31)
```

## Simple linear regression

Now we are finally ready to train a regression model.

Since the article is nominally about the attractiveness of the instructor, we will train the simple linear regression on the `bty_avg` feature.

In the cell that follows, write code to

- use `sklearn` to fit a simple linear regression model on the training set, using `bty_avg` as the feature on which to train. Save your fitted model in a variable `reg_simple`.
- print the intercept and coefficient of the model.
- use `predict` on the fitted model to estimate the evaluation score on the training set, and save this array in `y_pred_train`.
- use `predict` on the fitted model to estimate the evaluation score on the test set, and save this array in `y_pred_test`.

Then run the cell after that one, which will show you the training data, the test data, and your regression line.




```
reg_simple = LinearRegression().fit(train[['bty_avg']], train['score'])
reg_simple.coef_
reg_simple.intercept_
```

```
y_pred_train = reg_simple.predict(train[['bty_avg']])
y_pred_test = reg_simple.predict(test[['bty_avg']])
```

```
array([0.07367795])
```

```
3.842544653270749
```

```
sns.scatterplot(data=train, x="bty_avg", y="score", color='blue', alpha=0.5);
sns.scatterplot(data=test, x="bty_avg", y="score", color='green', alpha=0.5);
sns.lineplot(data=train, x="bty_avg", y=y_pred_train, color='red');
```



### Evaluate simple linear regression performance

Next, we will evaluate our model performance.

In the following cell, write a *function* to compute key performance metrics for your model:

- compute the R2 score on your training data
- compute the MSE on your training data
- compute the MSE, divided by the sample variance of `score`, on your training data. Recall that this metric tells us the ratio of average error of your model to average error of prediction by mean.
- and compute the same three metrics for your test set

```
def regression_performance(y_true_train, y_pred_train, y_true_test, y_pred_test):

    r2_train = metrics.r2_score(y_true_train, y_pred_train)
```

```

mse_train = metrics.mean_squared_error(y_true_train, y_pred_train)
norm_mse_train = metrics.mean_squared_error(y_true_train,
                                             y_pred_train)/(np.std(y_true_train)**2)

r2_test = metrics.r2_score(y_true_test, y_pred_test)
mse_test = metrics.mean_squared_error(y_true_test, y_pred_test)
norm_mse_test = metrics.mean_squared_error(y_true_test,
                                             y_pred_test)/(np.std(y_true_test)**2)

#print("Training:    %f %f %f" % (r2_train, mse_train, norm_mse_train))
#print("Test:       %f %f %f" % (r2_test, mse_test, norm_mse_test))

return [r2_train, mse_train, norm_mse_train, r2_test, mse_test, norm_mse_test]

```

Call your function to print the performance of the simple linear regression. Is a simple linear regression on `btv_avg` better than a “dumb” model that predicts the mean value of `score` for all samples?

```
vals = regression_performance(train['score'], y_pred_train, test['score'], y_pred_test)
```

## Multiple linear regression

Next, we’ll see if we can improve model performance using multiple linear regression, with more features included.

To start, we need to decide which features to use as input to our model. One possible approach is to use every feature in the dataset excluding the target variable, `score`.

You can build and view this list of features by running:

```

features = df_enc.columns.drop(['score'])
features

```

```

Index(['age', 'cls_perc_eval', 'cls_did_eval', 'cls_students', 'btv_f1lower',
      'btv_f1upper', 'btv_f2upper', 'btv_m1lower', 'btv_m1upper',
      'btv_m2upper', 'btv_avg', 'rank_teaching', 'rank_tenure track',
      'rank_tenured', 'ethnicity_minority', 'ethnicity_not minority',
      'gender_female', 'gender_male', 'language_english',
      'language_non-english', 'cls_level_lower', 'cls_level_upper',
      'cls_profs_multiple', 'cls_profs_single', 'cls_credits_multi credit',
      'cls_credits_one credit', 'pic_outfit_formal', 'pic_outfit_not formal',
      'pic_color_black&white', 'pic_color_color'],
      dtype='object')

```

In the following cell, write code to

- use `sklearn` to fit a linear regression model on the training set, using the `features` array as the list of features to train on. Save your fitted model in a variable `reg_multi`.
- print a table of the features used in the regression and the coefficient assigned to each. If you have saved your fitted regression in a variable named `reg_multi`, you can create and print this table with:

```

df_coef = pd.DataFrame(data =
                        {'feature': features,
                         'coefficient': reg_multi.coef_})
df_coef

```

```
reg_multi = LinearRegression().fit(train[features], train['score'])
df_coef = pd.DataFrame(data =
                        {'feature': features,
                         'coefficient': reg_multi.coef_})
df_coef
```

	feature	coefficient
0	age	-0.009493
1	cls_perc_eval	0.004385
2	cls_did_eval	0.002983
3	cls_students	-0.001427
4	btv_f1lower	7.022264
5	btv_f1upper	7.052930
6	btv_f2upper	7.032261
7	btv_m1lower	6.959419
8	btv_m1upper	6.999172
9	btv_m2upper	6.978719
10	btv_avg	-42.015575
11	rank_teaching	0.084572
12	rank_tenure track	-0.085993
13	rank_tenured	0.001421
14	ethnicity_minority	-0.131245
15	ethnicity_not minority	0.131245
16	gender_female	-0.121054
17	gender_male	0.121054
18	language_english	0.085742
19	language_non-english	-0.085742
20	cls_level_lower	-0.006088
21	cls_level_upper	0.006088
22	cls_profs_multiple	0.004518
23	cls_profs_single	-0.004518
24	cls_credits_multi credit	-0.283913
25	cls_credits_one credit	0.283913
26	pic_outfit_formal	0.039920
27	pic_outfit_not formal	-0.039920
28	pic_color_black&white	0.065689
29	pic_color_color	-0.065689

**Discussion Question 2** Look at the list of features and coefficients, especially those related to the attractiveness ratings.

Are these results surprising, based on the results of the simple linear regression? Explain your answer.

### Effect of collinearity

Note especially the coefficients associated with each of the individual attractiveness rankings, and the coefficient associated with the average attractiveness ranking. Each of these features separately seems to have a large effect; however, because they are strongly *collinear*, they cancel one another out.

(You should be able to see the collinearity clearly in the pairplot you created.)

In the following cell, write code to

- create a new `features` array, that drops the *individual* attractiveness rankings in addition to the `score` variable (but do *not* drop the average attractiveness ranking)
- use `sklearn` to fit a linear regression model on the training set, using the new `features` array as the list of features to train on. Save your fitted model in a variable `reg_avgbty`.
- print a table of the features used in the regression and the coefficient assigned to each.

```
features = df_enc.columns.drop(['score',
                                'bty_f1lower', 'bty_f1upper', 'bty_f2upper',
                                'bty_m1lower', 'bty_m1upper', 'bty_m2upper'])
reg_avgbty = LinearRegression().fit(train[features], train['score'])

df_coef = pd.DataFrame(data =
                        {'feature': features,
                         'coefficient': reg_avgbty.coef_})

df_coef
```

	feature	coefficient
0	age	-0.009297
1	cls_perc_eval	0.004900
2	cls_did_eval	0.003737
3	cls_students	-0.001740
4	bty_avg	0.040577
5	rank_teaching	0.075283
6	rank_tenure track	-0.059429
7	rank_tenured	-0.015854
8	ethnicity_minority	-0.111830
9	ethnicity_not minority	0.111830
10	gender_female	-0.097201
11	gender_male	0.097201
12	language_english	0.083435
13	language_non-english	-0.083435
14	cls_level_lower	-0.012437
15	cls_level_upper	0.012437
16	cls_profs_multiple	0.009897
17	cls_profs_single	-0.009897
18	cls_credits_multi credit	-0.285089
19	cls_credits_one credit	0.285089
20	pic_outfit_formal	0.053104
21	pic_outfit_not formal	-0.053104
22	pic_color_black&white	0.076615
23	pic_color_color	-0.076615

**Discussion Question 3** Given the model parameters you have found, which is associated with the strongest effect (on average) on the evaluation score:

- Instructor ethnicity
- Instructor gender

(Note that in general, we cannot use the coefficient to compare the effect of features that have a different range. But both ethnicity and gender are represented by binary one hot-encoded variables.)

## Evaluate multiple regression model performance

Evaluate the performance of your `reg_avgbty` model. In the next cell, write code to:

- use the `predict` function on your fitted regression to find  $\hat{y}$  for all samples in the *training* set, and save this in an array called `y_pred_train`
- use the `predict` function on your fitted regression to find  $\hat{y}$  for all samples in the *test* set, and save this in an array called `y_pred_test`
- call the `regression_performance` function you wrote in a previous cell, and print the performance metrics on the training and test set.

```
y_pred_train = reg_avgbty.predict(train[features])
y_pred_test = reg_avgbty.predict(test[features])

vals = regression_performance(train['score'], y_pred_train, test['score'], y_pred_test)
```

**Discussion Question 4** Based on the analysis above, what portion of the variation in instructor teaching evaluation can be explained by the factors unrelated to teaching performance, such as the physical characteristics of the instructor?

---

**Discussion Question 5** Based on the analysis above, is your model better at predicting instructor teaching scores than a “dumb” model that just assigns the mean teaching score to every instructor? Explain.

---

**Discussion Question 6** Suppose you are hired by the ECE department to develop a classifier that will identify high-performing faculty, who will then be awarded prizes for their efforts.

Based on the analysis above, do you think it would be fair to use scores on teaching evaluations as an input to your classifier? Explain your answer.

---

## Exploring unexpected correlation

There are some features that we do *not* expect to be correlated with the instructor’s score.

For example, consider the “features” related to the photograph used by the students who rated the instructor’s attractiveness.

There is no reason that characteristics of an instructor’s photograph - whether it was in black and white or color, how the instructor was dressed in the photograph - should influence the ratings of students in the instructor’s class. (These students did not even see the photograph.)

In the next cell, write code to

- create a new `features` array that drops the `score` variable, all of the individual attractiveness rankings, and the variables related to the photograph used for attractiveness rankings.
- use it to fit a model (saved in `reg_nopic`).
- use `reg_nopic` to predict the evaluation scores on both the training and test set
- compute the same set of metrics as above.

```
features = df_enc.columns.drop(['score',
                                'bty_f1lower', 'bty_f1upper', 'bty_f2upper',
                                'bty_m1lower', 'bty_m1upper', 'bty_m2upper',
                                'pic_outfit_formal', 'pic_outfit_not formal',
```

```

'pic_color_black&white', 'pic_color_color'])

reg_nopic = LinearRegression().fit(train[features], train['score'])

y_pred_train = reg_nopic.predict(train[features])
y_pred_test = reg_nopic.predict(test[features])

vals = regression_performance(train['score'], y_pred_train, test['score'], y_pred_test)

```

**Discussion Question 7** Is your model less predictive when features related to the instructor photograph are excluded? Explain.

When a machine learning model seems to use a feature that is not expected to be correlated with the target variable (such as the characteristics of the instructor's photograph...), this can sometimes be a signal that information is “leaking” between the training and test set.

In this dataset, each row represents a single course. However, some instructors teach more than one course, and an instructor might get similar evaluation scores on all of the courses he or she teaches.

(According to the paper for which this dataset was collected, 94 faculty members taught the 463 courses represented in the dataset, with some faculty members teaching as many as 13 courses.)

For example, consider the output of the following command, which prints all of the one credit courses in the data:

```
df.loc[df['cls_credits']=='one credit']
```

	score	rank	ethnicity	gender	language	age	\
124	3.5	teaching	not minority	female	english	52	
179	4.4	tenure track	minority	female	english	47	
185	4.6	tenure track	minority	female	english	47	
245	4.2	teaching	not minority	female	english	50	
246	4.7	teaching	not minority	female	english	50	
339	4.8	tenure track	not minority	male	english	43	
340	4.9	tenure track	not minority	male	english	43	
343	4.5	tenure track	not minority	male	english	43	
344	4.9	tenure track	not minority	male	english	43	
345	4.4	tenure track	not minority	male	english	43	
347	4.6	teaching	minority	male	english	50	
348	5.0	teaching	minority	male	english	50	
349	4.9	teaching	minority	male	english	50	
350	4.6	teaching	minority	male	english	50	
351	4.8	teaching	minority	male	english	50	
352	4.9	teaching	minority	male	english	50	
353	4.9	teaching	minority	male	english	50	
354	4.9	teaching	minority	male	english	50	
355	5.0	teaching	minority	male	english	50	
356	4.5	teaching	minority	male	english	50	
393	4.8	teaching	not minority	male	english	45	
394	4.2	teaching	not minority	male	english	45	
396	4.8	teaching	not minority	male	english	45	
409	4.7	teaching	not minority	female	english	47	

410	4.6	teaching	not minority	female	english	47
411	4.6	teaching	not minority	female	english	47
462	4.1	tenure track	minority	female	non-english	42

	cls_perc_eval	cls_did_eval	cls_students	cls_level	...	cls_credits	\
124	89.47369	17	19	upper	...	one credit	
179	100.00000	16	16	lower	...	one credit	
185	95.23810	20	21	lower	...	one credit	
245	75.00000	24	32	lower	...	one credit	
246	66.66666	14	21	lower	...	one credit	
339	53.57143	15	28	lower	...	one credit	
340	60.00000	18	30	lower	...	one credit	
343	94.44444	17	18	lower	...	one credit	
344	84.61539	22	26	lower	...	one credit	
345	60.00000	18	30	lower	...	one credit	
347	70.83334	17	24	lower	...	one credit	
348	90.90909	20	22	lower	...	one credit	
349	84.00000	21	25	lower	...	one credit	
350	88.46154	23	26	lower	...	one credit	
351	86.36364	19	22	lower	...	one credit	
352	76.92308	20	26	lower	...	one credit	
353	85.00000	17	20	lower	...	one credit	
354	81.81818	18	22	lower	...	one credit	
355	95.23810	20	21	lower	...	one credit	
356	90.47619	19	21	lower	...	one credit	
393	70.58823	12	17	lower	...	one credit	
394	85.00000	17	20	lower	...	one credit	
396	73.07692	19	26	lower	...	one credit	
409	88.23529	15	17	lower	...	one credit	
410	100.00000	10	10	lower	...	one credit	
411	94.11765	16	17	lower	...	one credit	
462	80.00000	28	35	lower	...	one credit	

	bty_f1lower	bty_f1upper	bty_f2upper	bty_m1lower	bty_m1upper	\
124	6	6	4	2	4	
179	2	6	6	3	5	
185	2	6	6	3	5	
245	2	3	5	2	3	
246	2	3	5	2	3	
339	3	4	4	2	4	
340	3	4	4	2	4	
343	3	4	4	2	4	
344	3	4	4	2	4	
345	3	4	4	2	4	
347	1	5	4	1	4	
348	1	5	4	1	4	
349	1	5	4	1	4	
350	1	5	4	1	4	
351	1	5	4	1	4	
352	1	5	4	1	4	
353	1	5	4	1	4	
354	1	5	4	1	4	
355	1	5	4	1	4	
356	1	5	4	1	4	

393	1	4	2	5	4
394	1	4	2	5	4
396	1	4	2	5	4
409	8	6	6	4	9
410	8	6	6	4	9
411	8	6	6	4	9
462	3	8	7	4	6

	bty_m2upper	bty_avg	pic_outfit	pic_color
124	7	4.833	not formal	color
179	4	4.333	not formal	color
185	4	4.333	not formal	color
245	4	3.167	not formal	color
246	4	3.167	not formal	color
339	4	3.500	not formal	color
340	4	3.500	not formal	color
343	4	3.500	not formal	color
344	4	3.500	not formal	color
345	4	3.500	not formal	color
347	5	3.333	not formal	color
348	5	3.333	not formal	color
349	5	3.333	not formal	color
350	5	3.333	not formal	color
351	5	3.333	not formal	color
352	5	3.333	not formal	color
353	5	3.333	not formal	color
354	5	3.333	not formal	color
355	5	3.333	not formal	color
356	5	3.333	not formal	color
393	4	3.333	not formal	color
394	4	3.333	not formal	color
396	4	3.333	not formal	color
409	7	6.667	not formal	black&white
410	7	6.667	not formal	black&white
411	7	6.667	not formal	black&white
462	4	5.333	not formal	color

[27 rows x 21 columns]

We observe that 10 out of 27 one-credit courses are taught by what seems to be the same instructor - we don't know his name, but let's call him John. John is a teaching-track professor, minority ethnicity, male, English-language trained, 50 years old, average attractiveness 3.333, and whose photograph is in color and not formal.

This provides a clue regarding the apparent importance of the `cls_credits` variable and other “unexpected” variables in predicting the teaching score.

Certain variables may be used by the model to identify the instructor, and then learn a relationship between the *individual instructor* and his or her typical evaluation score, instead of learning a true relationship between the *variable* and the evaluation score.

To see if this is plausible, let's add an “instructor ID” to each row in our data frame. The data set doesn't include an instructor ID column, but we can still uniquely identify every instructor by looking at the combination of rank, ethnicity, gender, language of training, age, attractiveness score, and characteristics of the photo (formal or not, black and white or color).



```

instructor_id = df[['rank', 'ethnicity', 'gender', 'language',
                    'pic_outfit', 'pic_color']].agg('-'.join, axis=1)
instructor_id += ' -' + df['age'].astype(str)
instructor_id += ' -' + df['bty_avg'].astype(str)

df_enc = df_enc.assign(instructor_id = instructor_id)

df_enc['instructor_id'].head()

```

```

0    tenure track-minority-female-english-not forma...
1    tenure track-minority-female-english-not forma...
2    tenure track-minority-female-english-not forma...
3    tenure track-minority-female-english-not forma...
4    tenured-not minority-male-english-not formal-c...
Name: instructor_id, dtype: object

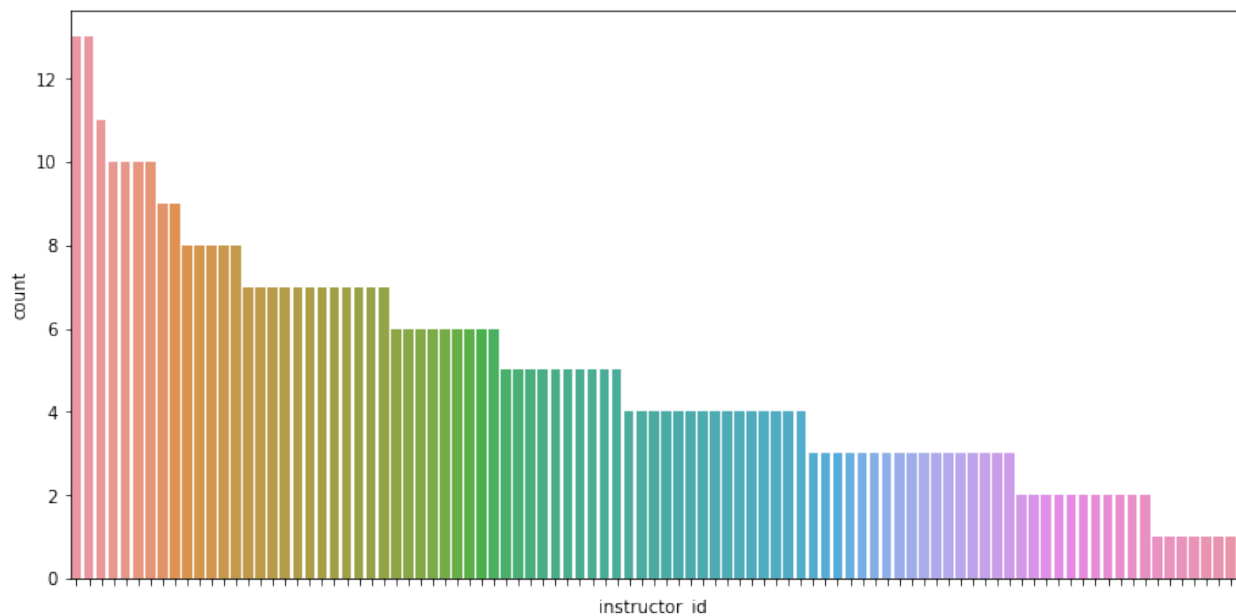
```

Let's plot the frequency with which each "instructor ID" appears in the data:

```

_ = plt.figure(figsize=(12,6))
ax = sns.countplot(x="instructor_id", data=df_enc,
                  order=df_enc['instructor_id'].value_counts().index)
_ = ax.set(xticklabels=[])

```



There are 95 unique instructor IDs. According to the paper, the data set includes 94 instructors. We are working with a slightly modified version of the data from the paper. It seems we've been able to uniquely identify every instructor.

Some instructors are represented as many as 13 times in the dataset. Only a handful of instructors appear only once in the data.


Furthermore, we can see that most instructors get similar scores for all of the courses they teach, with a few exceptions:

```

score_order =
    df_enc.groupby('instructor_id')['score'].agg('mean').sort_values(ascending=False).index

_ = plt.figure(figsize=(12,6))
ax = sns.boxplot(x=df_enc['instructor_id'], y=df_enc['score'],
                order=score_order,
                color='white', width=0.4)
ax = sns.stripplot(x=df_enc['instructor_id'], y=df_enc['score'],
                  order=score_order)
_ = ax.set(xticklabels=[])

```



To explore this issue further, we will repeat our analysis using two different ways of splitting the dataset:

1. **Random split:** shuffle data and split it into training and test sets. Train the model using the training data, then evaluate its performance on the test set. (This is what we have done so far.)
2. **Group split:** split data into training and test sets in a way that ensures that each individual *instructor* is represented in either the training data or the test data, but not both. Train the model using the training data, then evaluate its performance on the test set. If the model is “memorizing” individual instructors, rather than learning a general relationship between features and teaching evaluation score, it will have much worse performance on the test set, because it has to predict scores for instructors it hasn’t “seen” yet.

Because the dataset is small, the performance evaluation may be influenced by the random sample of rows that happen to end up in the training vs. test set. (If a few rows more rows than usual that are very “easy” to predict are placed in the test set, we might see better performance than we would with a different test set.) So, we will also repeat the splitting procedure several times, and look at the *average* performance across different train-test splits.

```

n_splits = 10
metrics_rs = np.zeros((n_splits, 6))
rs = model_selection.KFold(n_splits=n_splits, shuffle=True)

for i, split in enumerate(rs.split(df_enc)):
    train_idx, test_idx = split
    train = df_enc.iloc[train_idx]

```

```

test = df_enc.iloc[test_idx]

features = df_enc.columns.drop(['score', 'instructor_id'])

# train a multiple linear regression using
# the train dataset and the list of features created above
# save the fitted model in reg_rndsplit
# then use the model to create y_pred_train and y_pred_test,
# the model predictions on the training set and test set.
# Finally, return the array of model performance metrics

reg_rndsplit = LinearRegression().fit(train[features], train['score'])

y_pred_train = reg_rndsplit.predict(train[features])
y_pred_test = reg_rndsplit.predict(test[features])

metrics_rs[i] = regression_performance(train['score'], y_pred_train, test['score'],
                                       y_pred_test)

np.mean(metrics_rs, axis=0)

array([0.21665477, 0.23113384, 0.78334523, 0.10256849, 0.25931647,
       0.89743151])


```

With this approach to splitting the data, the model appears to have some predictive value on the test set (which is supposed to represent performance on “new” data.)

```

_ = plt.bar(x=['Train', 'Test'], height=np.mean(metrics_rs, axis=0)[[0,3]])
_ = plt.ylabel("R2")
_ = plt.ylim(0, 1)

```



Next, we will perform our splits, train a model, and get performance metrics according to the second scheme, in which an instructor may be present in either the training set or the test set, but not both.

```
n_splits = 10
metrics_gs = np.zeros((n_splits, 6))
gs = model_selection.GroupKFold(n_splits=n_splits)

for i, split in enumerate(gs.split(df_enc,
                                   df_enc['score'],
                                   df_enc['instructor_id'])):

    train_idx, test_idx = split

    train = df_enc.iloc[train_idx]
    test = df_enc.iloc[test_idx]

    features = df_enc.columns.drop(['score', 'instructor_id'])

    # train a multiple linear regression using
    # the train dataset and the list of features created above
    # save the fitted model in reg_grpsplit
    # then use the model to create y_pred_train and y_pred_test,
    # the model predictions on the training set and test set.
    # Finally, return the array of model performance metrics

    reg_grpsplit = LinearRegression().fit(train[features], train['score'])

    y_pred_train = reg_grpsplit.predict(train[features])
    y_pred_test = reg_grpsplit.predict(test[features])

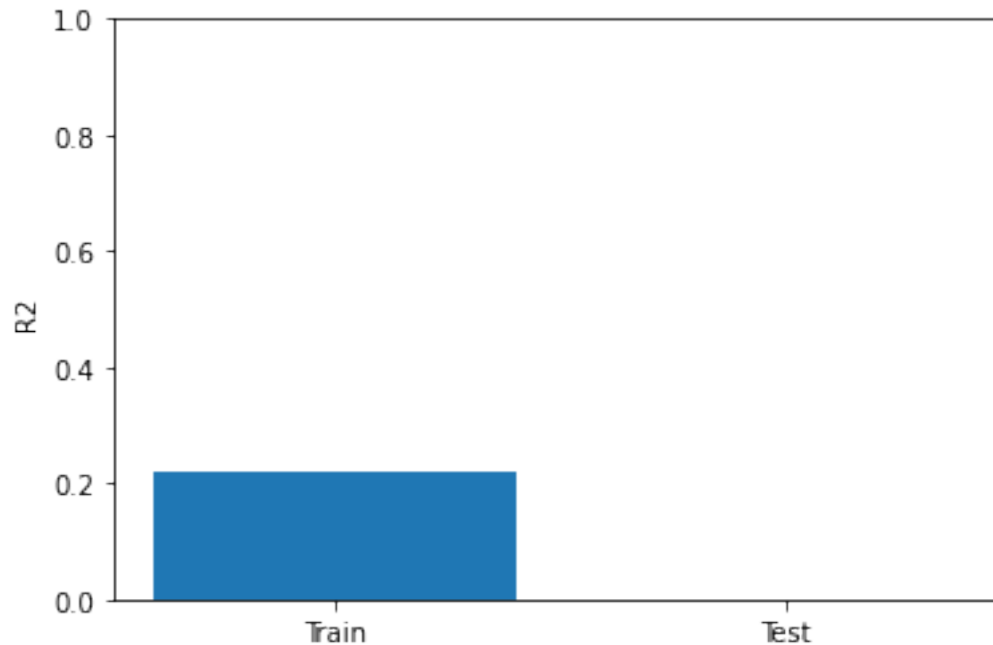
    metrics_gs[i] = regression_performance(train['score'], y_pred_train, test['score'],
                                           y_pred_test)

np.mean(metrics_gs, axis=0)

array([ 0.22092182,  0.22984319,  0.77907818, -0.03138834,  0.29278308,
        1.03138834])
```

With the second approach to splitting the data, the model has no predictive value on the test set.

```
_ = plt.bar(x=['Train', 'Test'], height=np.mean(metrics_gs, axis=0)[[0,3]])
_ = plt.ylabel("R2")
_ = plt.ylim(0, 1)
```



**Discussion Question 8** Based on your analysis above, do you think your model will be useful to predict the teaching evaluation scores of a new faculty member at UT Austin, based on his or her physical characteristics and the characteristics of the course?

---

### Review: what went wrong?

In this case study, we saw *two* problems:

The first problem is that the model was “memorizing” the individual instructors that appeared in the training data, rather than learning a general relationship between the features and the target variable. This is known as *overfitting*.

Usually, when a model is overfitting, it will be evident in the evaluation on the test set, because a model that overfits on training data will have excellent performance on training data and poor performance on test data. That’s where the second problem comes in: data leakage! We expect the model to be able to predict a baseline score for instructors it has not been trained on, but our model was being trained on data from a set of instructors, then evaluated on data from the same instructors.

As a result of this data leakage, the model had overly optimistic error on the test set. The model appeared to generalize to new, unseen, data, but in fact would not generalize to different instructors.

One of the “red flags” that helped us identify the problem was that the model seemed to be learning from features that we know are not really informative - for example, the characteristics of the photo used to derive the attractiveness ratings. This is often a sign of data leakage.